

Part 8 :

RHADOOP

R:

R is an open-source programming language that is extensively used for **statistical and graphical analysis**.

One major quality of **R's** is that it produces well-designed **quality plots** with greater ease, including mathematical **symbols and formulae where needed**.


Some reasons for which R is considered the best fit for data analytics :

- A robust collection of **packages**
- Powerful **data visualization** techniques
- Commendable **Statistical and graphical** programming features
- Object-oriented programming language
- It has a wide smart **collection of operators** for calculations of arrays, particular matrices, etc
- Graphical representation capability on display or on hard copy.

R:

R is very powerful for statistical analysis, However:

- ✓ All the calculations are performed by loading entire data in RAM
- ✓ Scaling up RAM has an upper limit
- ✓ Hadoop framework allows parallel processing of massive amount of data
- ✓ Using R with Hadoop facilitates horizontal scalability of statistical calculations

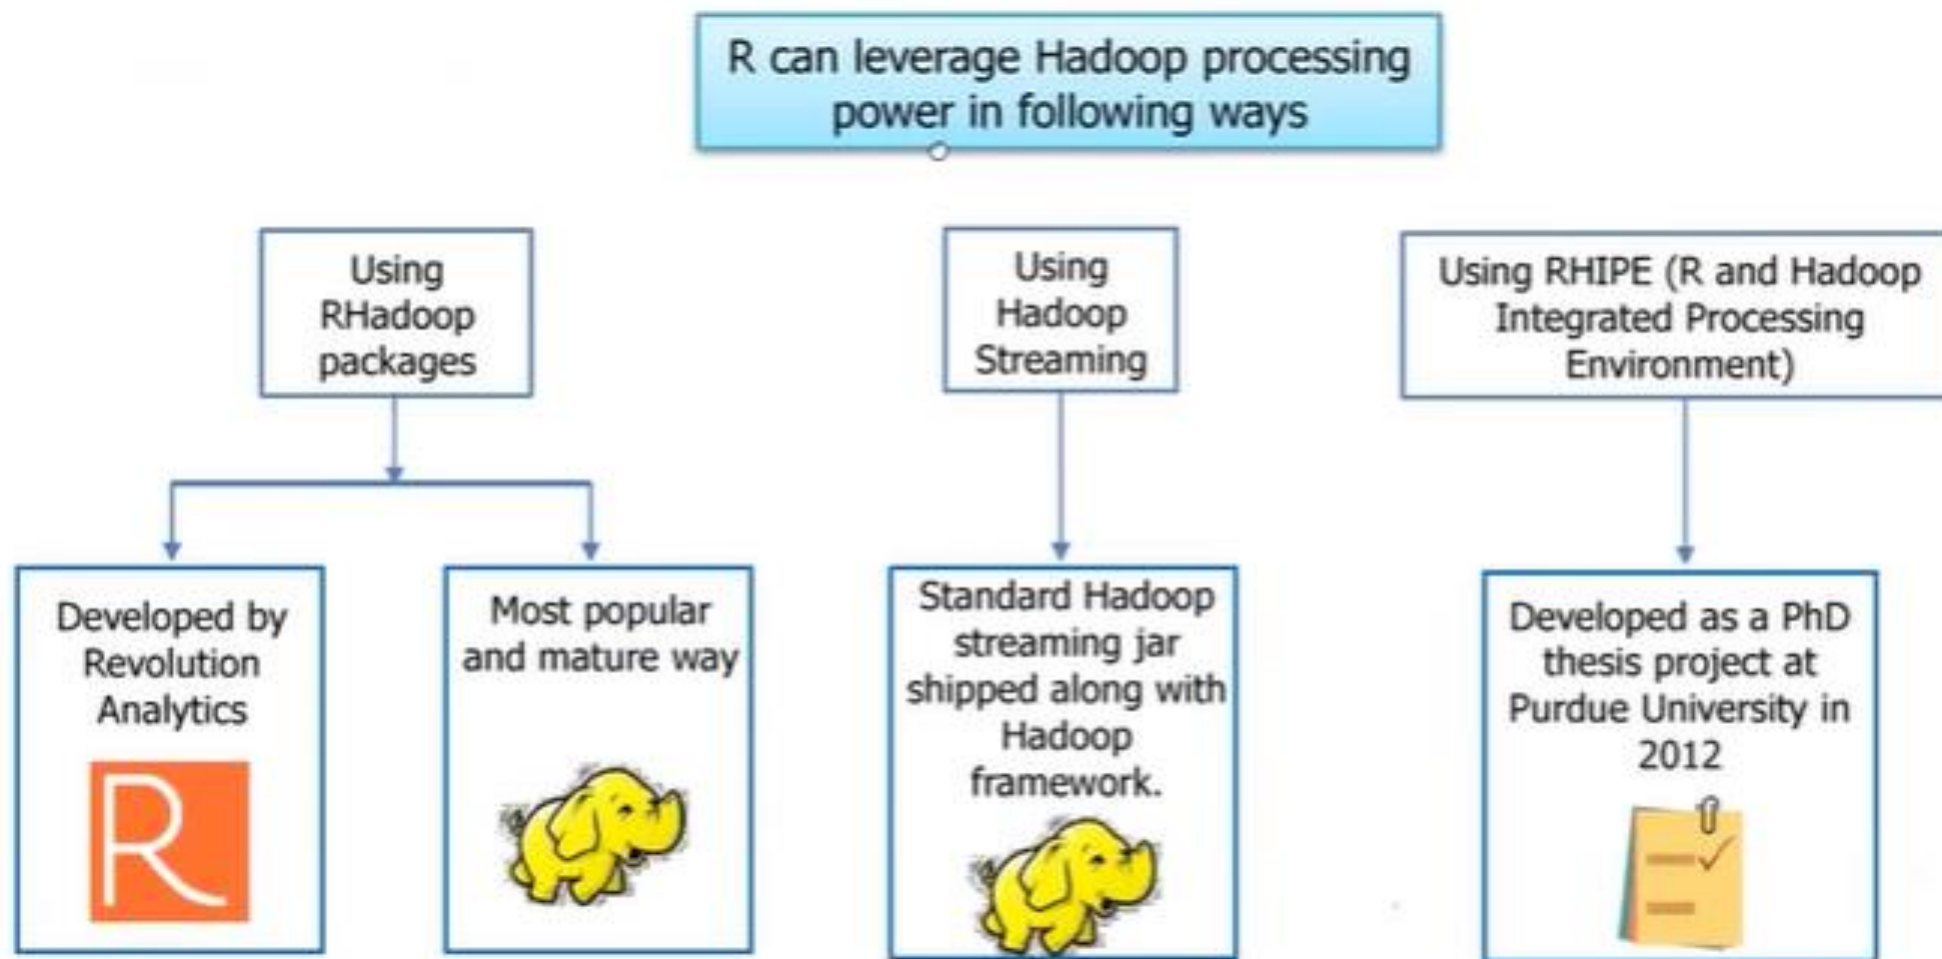


- **Rhadoop:** Collection of 3 packages for performing large data operations with an R environment.

Hadoop & R

- Hadoop for R
 - Datasets that do not fit in memory but can be naturally partitioned
 - Harness well known infrastructure for embarrassingly parallel tasks
- R for Hadoop
 - Very large number of implementations of different functions and libraries (machine learning, text processing etc)
 - Native C implementation – very fast number crunching, matrix processing
 - All in-memory calculations


Integrate R and Hadoop: Different Ways



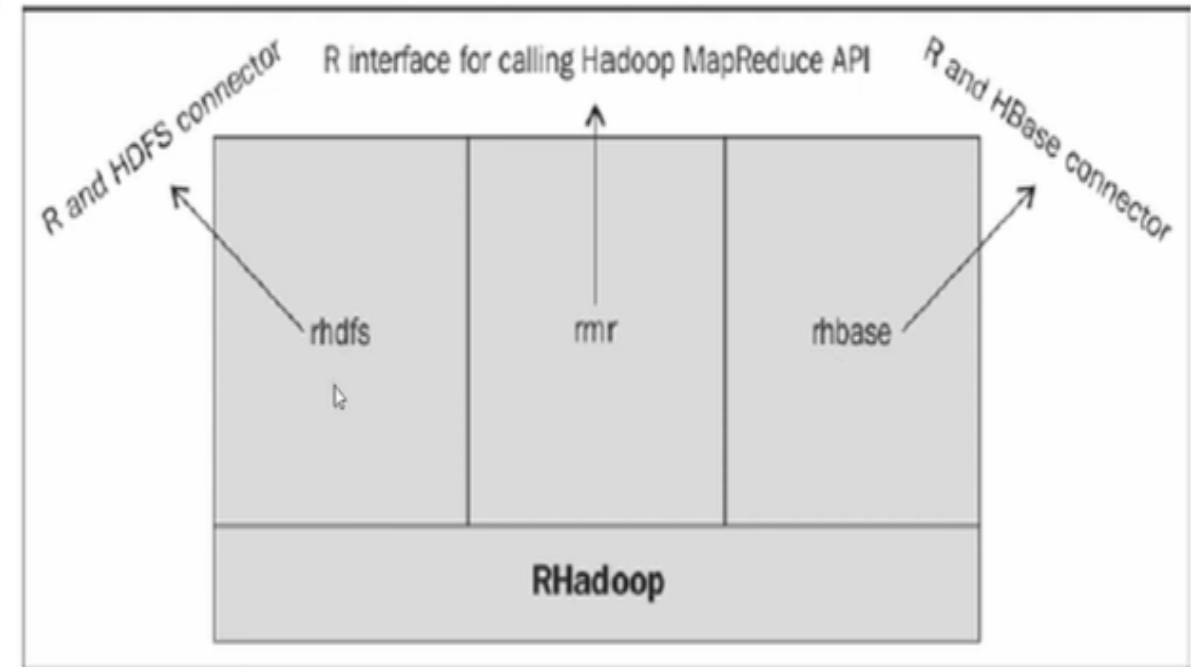
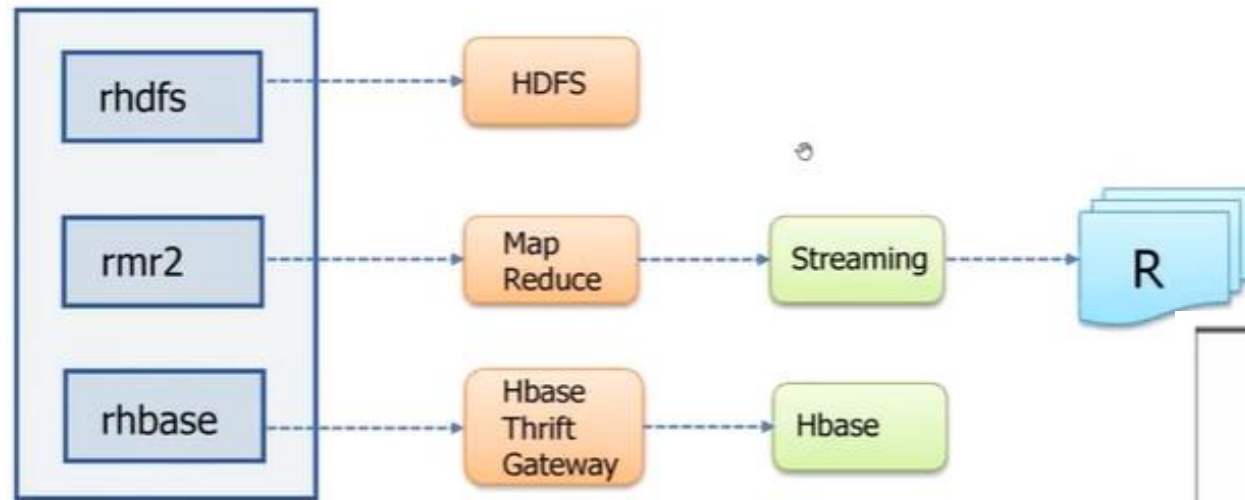
Exploring RHadoop

- ✓ Development lead by Revolution Analytics. Available on github: <https://github.com/RevolutionAnalytics/RHadoop>
- ✓ Mostly implemented in native R. Some pieces written in C++ and Java
- ✓ Contains the following packages:
 - **rmr**
 - Interface for running map/reduce jobs via Hadoop Streaming in R
 - **rhbase**
 - Interface to read/write data to/from HBase tables
 - **rhdfs**
 - Provides file manipulation capabilities for HDFS

- **rmr**: Helps to execute the Mapping & Reducing codes in R
- **rhdfs**: Provides file management capabilities by integrating with HDFS.
- **rhbase**: provides R database management capability with integration with Hbase.



RHadoop Architecture



RHadoop Ecosystem

rmr2



rmr2

Enables writing MapReduce jobs using R

Ability to parallelize algorithms

Ability to use big data sets without needing to sample data

Mapreduce(input, output, map, reduce, ...)

Reduces takes a key and a collection of values which could be vector, list, data frame or matrix

2.2.1 adds support for Windows (using HDP)

Functionalities of rmr:

- ✓ **Convenience**

- `keyval()`: used in generating output from input formatters, mappers, and reducers

- ✓ **Input/Output**

- `to.dfs()`
- `from.dfs()`
- `make.input.format()`
- `make.output.format()`

- ✓ **Job Execution**

- `mapreduce()`: submits job, returns a HDFS path pointing to output data



rmr Advantages:

✓ Well designed API

- User code needs only the manipulation of R objects (data frames, lists, vectors, strings, etc.)

✓ Flexible I/O subsystem

- Supports common input formats like CSV
- Provides control for input parsing (buffered, line-by-line w/o having to interact with stdin and stdout directly)

✓ Capability to construct chains of mapreduce jobs

- Result of mapreduce() is simply the HDFS path of job's output

✓ Local backend for quick prototyping

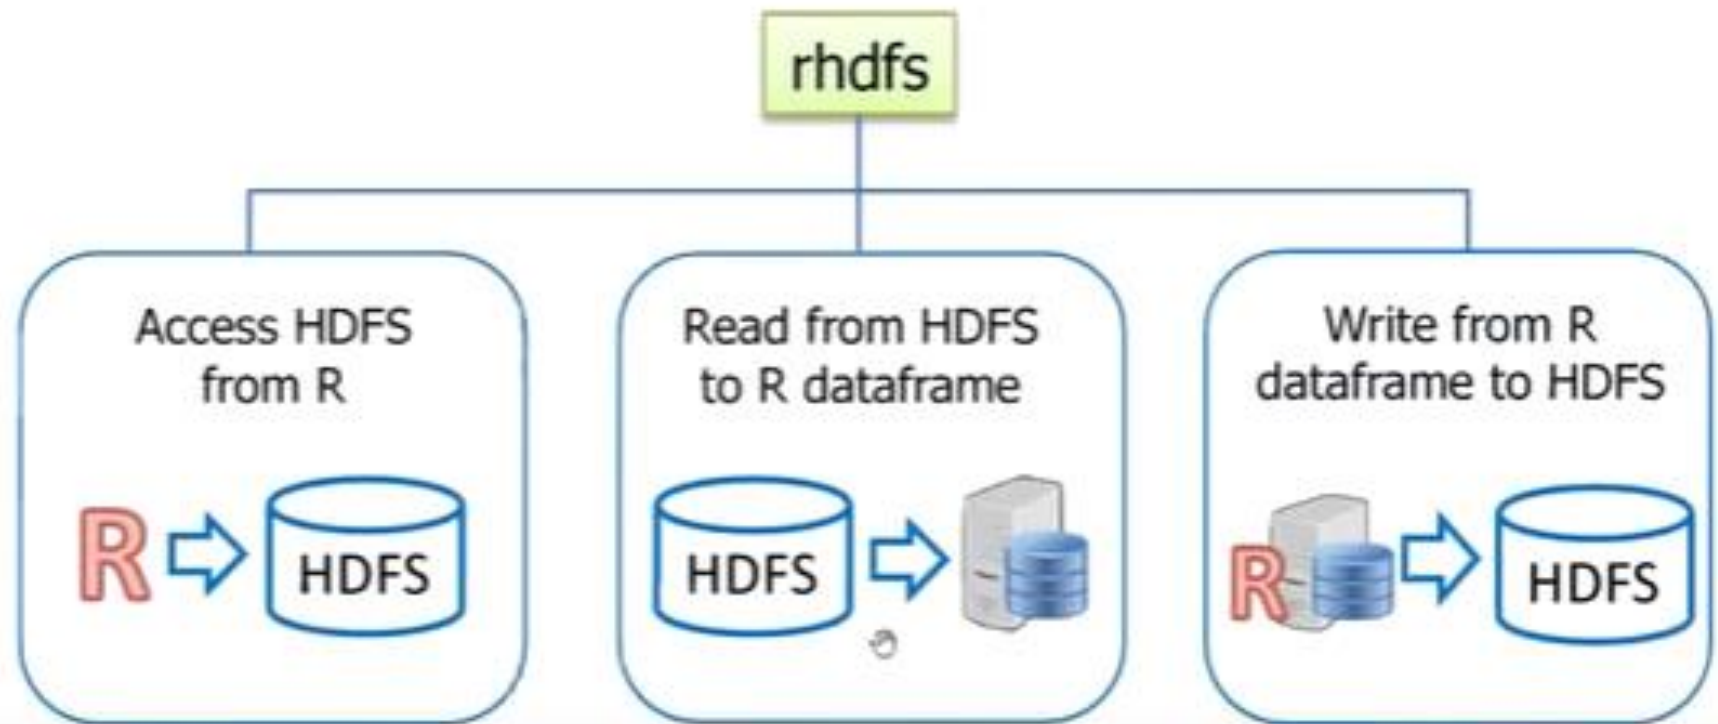
- The "local backend" refers to **running MapReduce jobs** on a **single machine**, utilizing the **local file system** instead of Hadoop's distributed file system (HDFS) and the **local computing resources** instead of a Hadoop cluster. This setup is primarily used for **development, testing, and prototyping purposes**.

- It provides a **convenient interface for R users** to write **MapReduce** programs and execute them on Hadoop clusters **without needing to write Java code directly**.

- users have **full control over how input data is processed and transformed into key-value pairs** for further processing.
- Users can define **custom output formats** or select from predefined formats like text or sequence file.

- MapReduce job will be executed **locally** within the R environment on the same machine where R is running.

rhdfs



rhdfs

- ✓ Hadoop CLI Commands & rhdfs equivalent
- ✓ `hadoop fs -ls /`
 - `hdfs.ls("/")`
- ✓ `hadoop fs -mkdir /user/rhdfs/ppt`
 - `hdfs.mkdir("/user/rhdfs/ppt")`
- ✓ `hadoop fs -put 1.txt /user/rhdfs/ppt/`
 - `localData <- system.file(file.path("unitTestData", "1.txt"), package="rhdfs")`
 - `hdfs.put(localData, "/user/rhdfs/ppt/1.txt")`
- ✓ `hadoop fs -get /user/rhdfs/ppt/1.txt 1.txt`
 - `hdfs.get("/user/rhdfs/ppt/1.txt", "test")`
- ✓ `hadoop fs -rm /user/rhdfs/ppt/1.txt`
 - `hdfs.delete("/user/rhdfs/ppt/1.txt")`

Functionalities in HDFS

- ✓ **File manipulations**
 - `hdfs.copy`, `hdfs.move`, `hdfs.rename`,
 - `hdfs.delete`, `hdfs.rm`, `hdfs.del`,
 - `hdfs.chown`, `hdfs.put`, `hdfs.get`
- ✓ **File read/write**
 - `hdfs.file`, `hdfs.write`, `hdfs.read`,
 - `hdfs.close`, `hdfs.flush`, `hdfs.seek`, `hdfs.tell`,
 - `hdfs.line.reader`, `hdfs.read.text.file`
- ✓ **Directory**
 - `hdfs.dircreate`, `hdfs.mkdir`
- ✓ **Utility**
 - `hdfs.ls`, `hdfs.list.files`, `hdfs.file.info`, `hdfs.exists`
- ✓ **Initialization**
 - `hdfs.init` and `hdfs.defaults`

`hdfs.copy, hdfs.move, hdfs.rename, hdfs.delete, hdfs.rm, hdfs.del,`
`hdfs.chown, hdfs.put, hdfs.get`

- File Read/Write

`hdfs.file, hdfs.write, hdfs.close, hdfs.flush, hdfs.read, hdfs.seek,`
`hdfs.tell, hdfs.line.reader, hdfs.read.text.file`

- Directory

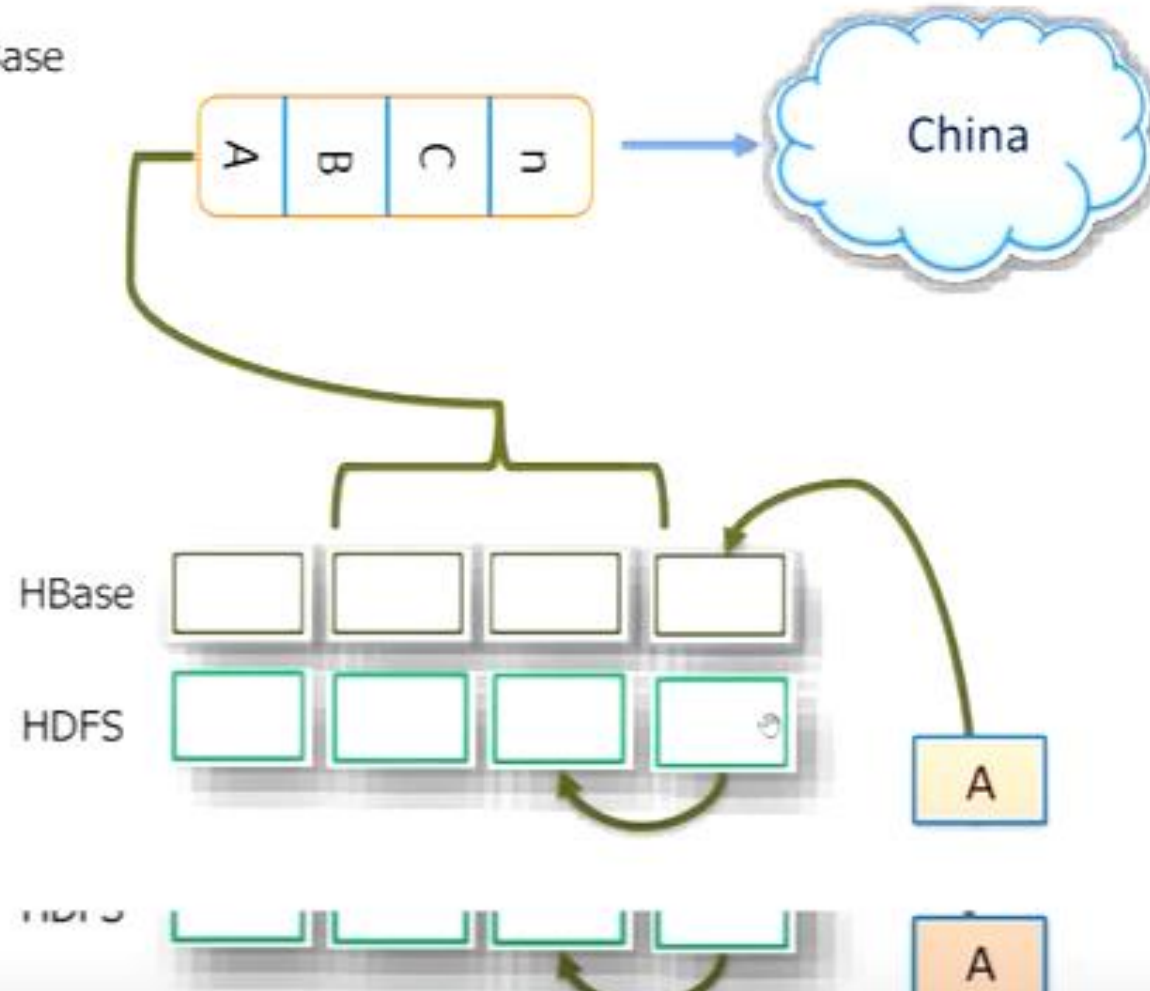
`hdfs.dircreate, hdfs.mkdir`

- Utility

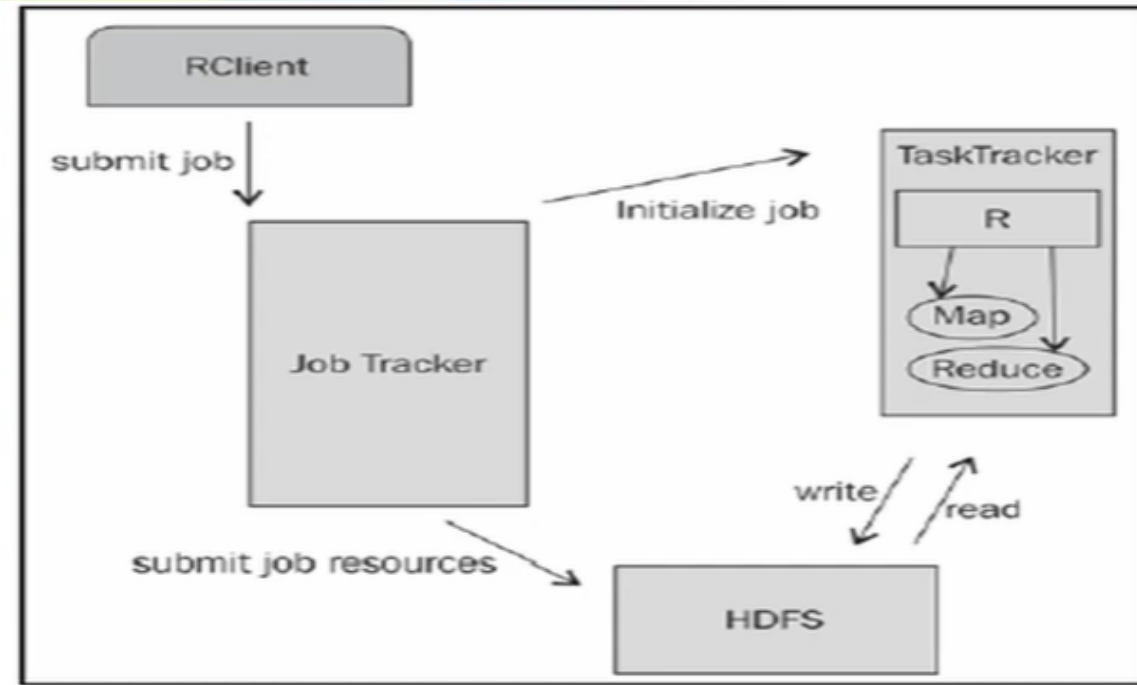
`hdfs.ls, hdfs.list.files, hdfs.file.info, hdfs.exists`

rhbase

- ✓ Access and change data within HBase
- ✓ Uses Thrift API
- ✓ Command Examples
 - `hb.new.table`
 - `hb.insert`
 - `hb.scan.ex`
 - `hb.scan`



RHIPE:



Components of RHIPE

- RHIPE: Stands for **R and Hadoop Integrated Programming Environment**
- RHIPE involves working with R and Hadoop Integrated Programming Environment
- **R client** : submits Query (Job) to job tracker.
- **Job Tracker** : Main coordinator – Responsible to allocating job/assign task/job to task tracker.
- **Task Tracker** : May be present in same location or different nodes.
- **Task Tracker** : **Perform Map Reduce** - Responsible for “Processing of Data”.
- Parallel to distributed computing.
- **HDFS**: **During Processing** : It can read and write data from HDFS

RHIPE Components:

RClient:

RClient is an R application that calls the JobTracker to execute the job with an indication of several MapReduce job resources such as Mapper, Reducer, input format, output format, input file, output file, and other several parameters that can handle the MapReduce jobs with RClient.

• JobTracker:

A JobTracker is the master node of the Hadoop Map Reduce operations for initializing and monitoring the MapReduce jobs over the Hadoop cluster..

RHIPE Components:

Task Tracker:

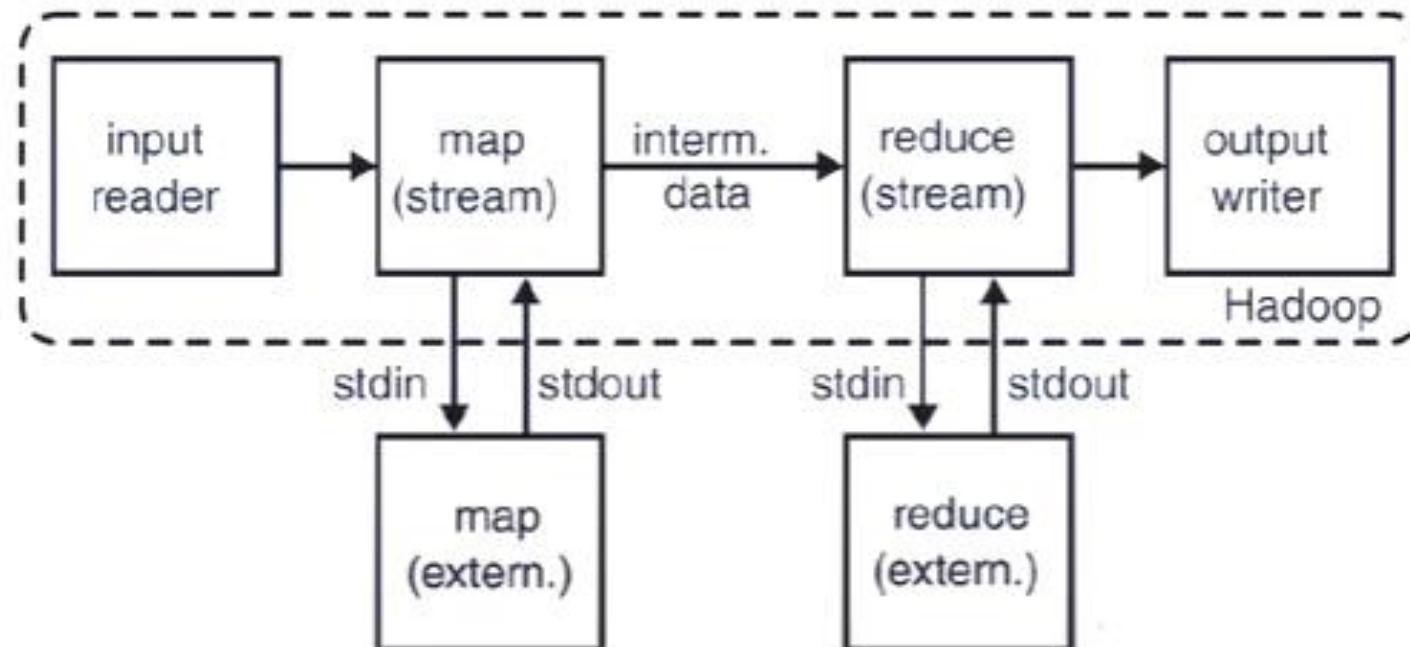
Task Tracker is a slave node in the Hadoop cluster. It executes the Map Reduce jobs as per the orders given by Job Tracker, retrieve the input data chunks, and run R-specific Mapper and Reducer over it. Finally, the output will be written on the HDFS directory.

• HDFS:

HDFS is a file system distributed over Hadoop clusters with several data nodes. It provides data services for various data operations.

Hadoop Streaming:

- ✓ A utility that allows one to run map/reduce jobs on the cluster with any executable; e.g. shell scripts, Perl, Python, etc
- ✓ Executable reads input from STDIN and writes output to STDOUT
- ✓ Maps input line data from STDIN to (key, value) pairs in STDOUT, in our case





ORCH:

- ORCH is known as Oracle R Connector.
- It helps in accessing the Hadoop cluster with the help of R.
- It allows us to manipulate the data residing in the Hadoop Distributed File System.

MINING DATA STREAMS

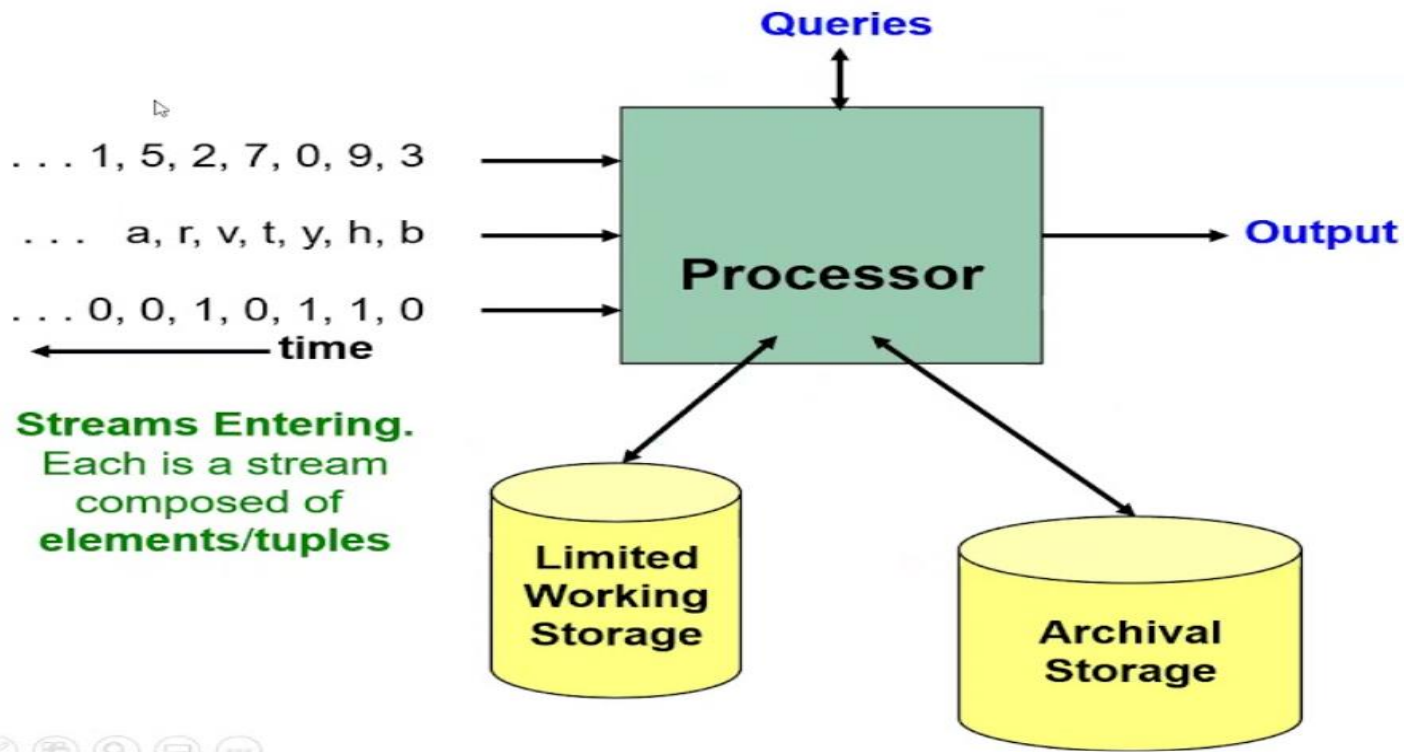
Data Stream ?

- Stream in English “a continuous flow”
- In computing
 - transmit or receive (data, especially video and audio material) over the Internet as a steady, continuous flow.

DBMS versus DSMS

- | | |
|-----------------------------------------------------------------|-----------------------------------------------------------|
| ■ One-time queries | ■ Continuous queries |
| ■ Random access | ■ Sequential access |
| ■ “Unbounded” disk store | ■ Bounded main memory |
| ■ Only current state matters | ■ Historical data is important |
| ■ No real-time services | ■ Real-time requirements |
| ■ Relatively low update rate | ■ Possibly multi-GB arrival rate |
| ■ Data at any granularity | ■ Data at fine granularity |
| ■ Assume precise data | ■ Data stale/imprecise |
| ■ Access plan determined by query processor, physical DB design | ■ Unpredictable/variable data arrival and characteristics |

General Stream Processing Model



Challenges in Working with Streaming Data

- A storage layer
- A processing layer

Needs plan for scalability, data durability, and fault tolerance in both the storage and processing layers.

Stream Processing

Parameter	Stream processing
Data scope	Queries or processing over data within a rolling time window, or on just the most recent data record.
Data size	Individual records or micro batches consisting of a few records.
Performance	Requires latency in the order of seconds or milliseconds.
Analyses	Simple response functions, aggregates, and rolling metrics.

Streaming Data Examples

- Sensors in transportation vehicles, industrial equipment, and farm machinery.
- A financial institution
- A real-estate website
- A solar power company
- A media publisher
- An online gaming company

Types of queries one wants answer on a data stream



Applications (1)

– Filtering a data stream

- Select elements with property x from the stream

– Counting distinct elements

- Number of distinct elements in the last k elements of the stream

– Estimating moments

- Estimate avg./std. dev. of last k elements

– Finding frequent elements

• Mining query streams

- Google wants to know what queries are more frequent today than yesterday

• Mining click streams

- Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour

• Mining social network news feeds

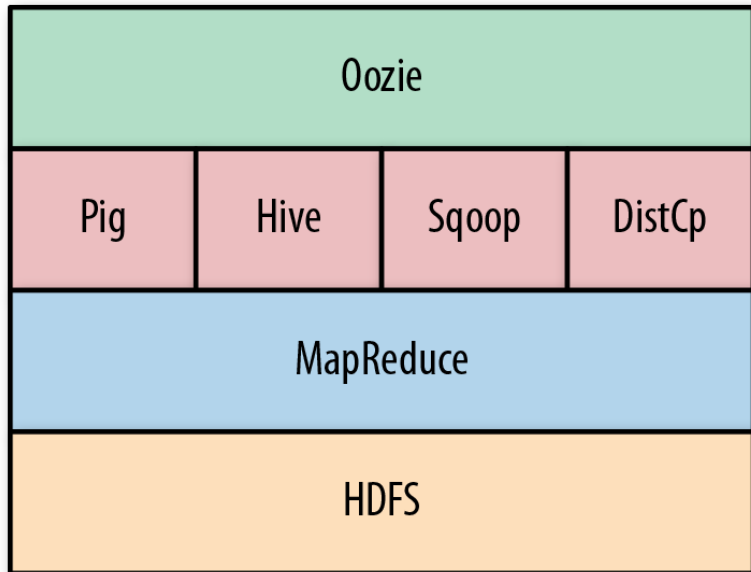
- E.g., look for trending topics on Twitter, Facebook

STREAM DATA PROCESSING METHODS

- ❖ **Random sampling** (but without knowing the total length in advance)
 - ✓ **Reservoir sampling**: maintain a set of s candidates in the reservoir, which form a true random sample of the element seen so far in the stream. As the data stream flows, every new element has a certain probability (s/N) of replacing an old element in the reservoir.
- ❖ **Sliding windows**
 - ✓ Make decisions based only on **recent** data of sliding window size w
 - ✓ An element arriving at time t expires at time $t + w$
- ❖ **Histograms**
 - ✓ Approximate the frequency distribution of element values in a stream
 - ✓ Partition data into a set of contiguous buckets
 - ✓ Equal-width (equal value range for buckets) vs. V-optimal (minimizing frequency variance within each bucket)

Apache OOOIE

- Apache Oozie is a workflow director designed to run and manage multiple related Apache Jobs.
- For instance, complete data input and analysis may require several related Hadoop jobs to be run as a workflow in which the output of one job may input to the successive job.
- Oozie is designed to construct and manage these workflows.
- Oozie is not a substitute for YARN scheduler.
- YARN manages the resources of individual jobs and Oozie provides a mechanism to connect and control Hadoop jobs on



Apache OOOIE

Oozie workflow jobs are represented using **Directed Acyclic Graphs**.

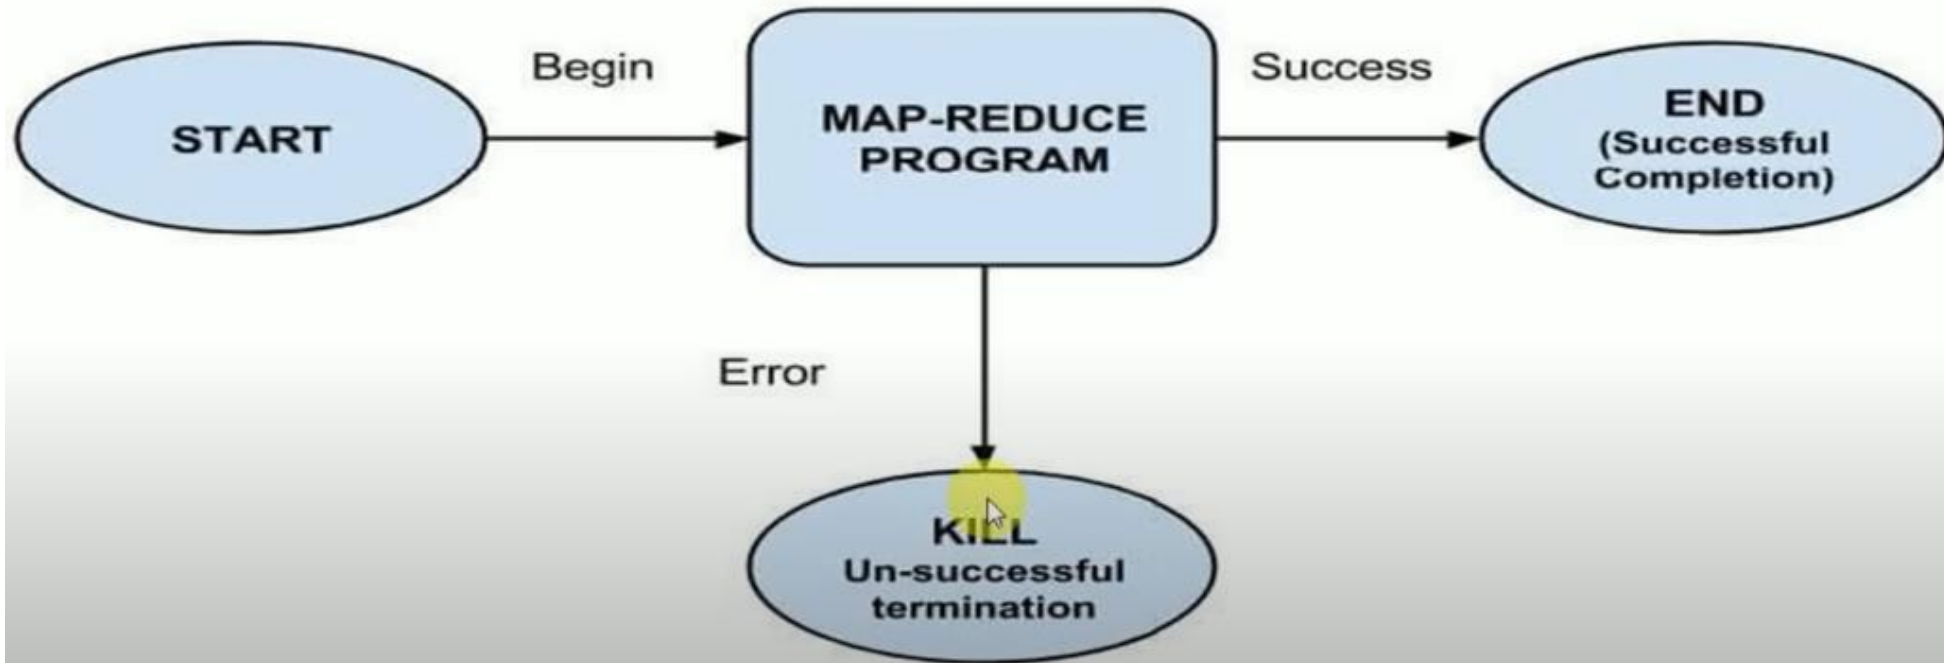
Oozie Supports three types of Jobs:

- **Workflow engine:** Responsibility of a workflow engine is to store and run workflows composed of Hadoop jobs e.g., MapReduce, Pig, Hive.
- **Coordinator engine:** It runs workflow jobs based on predefined schedules and availability of data.
- **Bundle:** Higher level abstraction that will batch a set of coordinator jobs

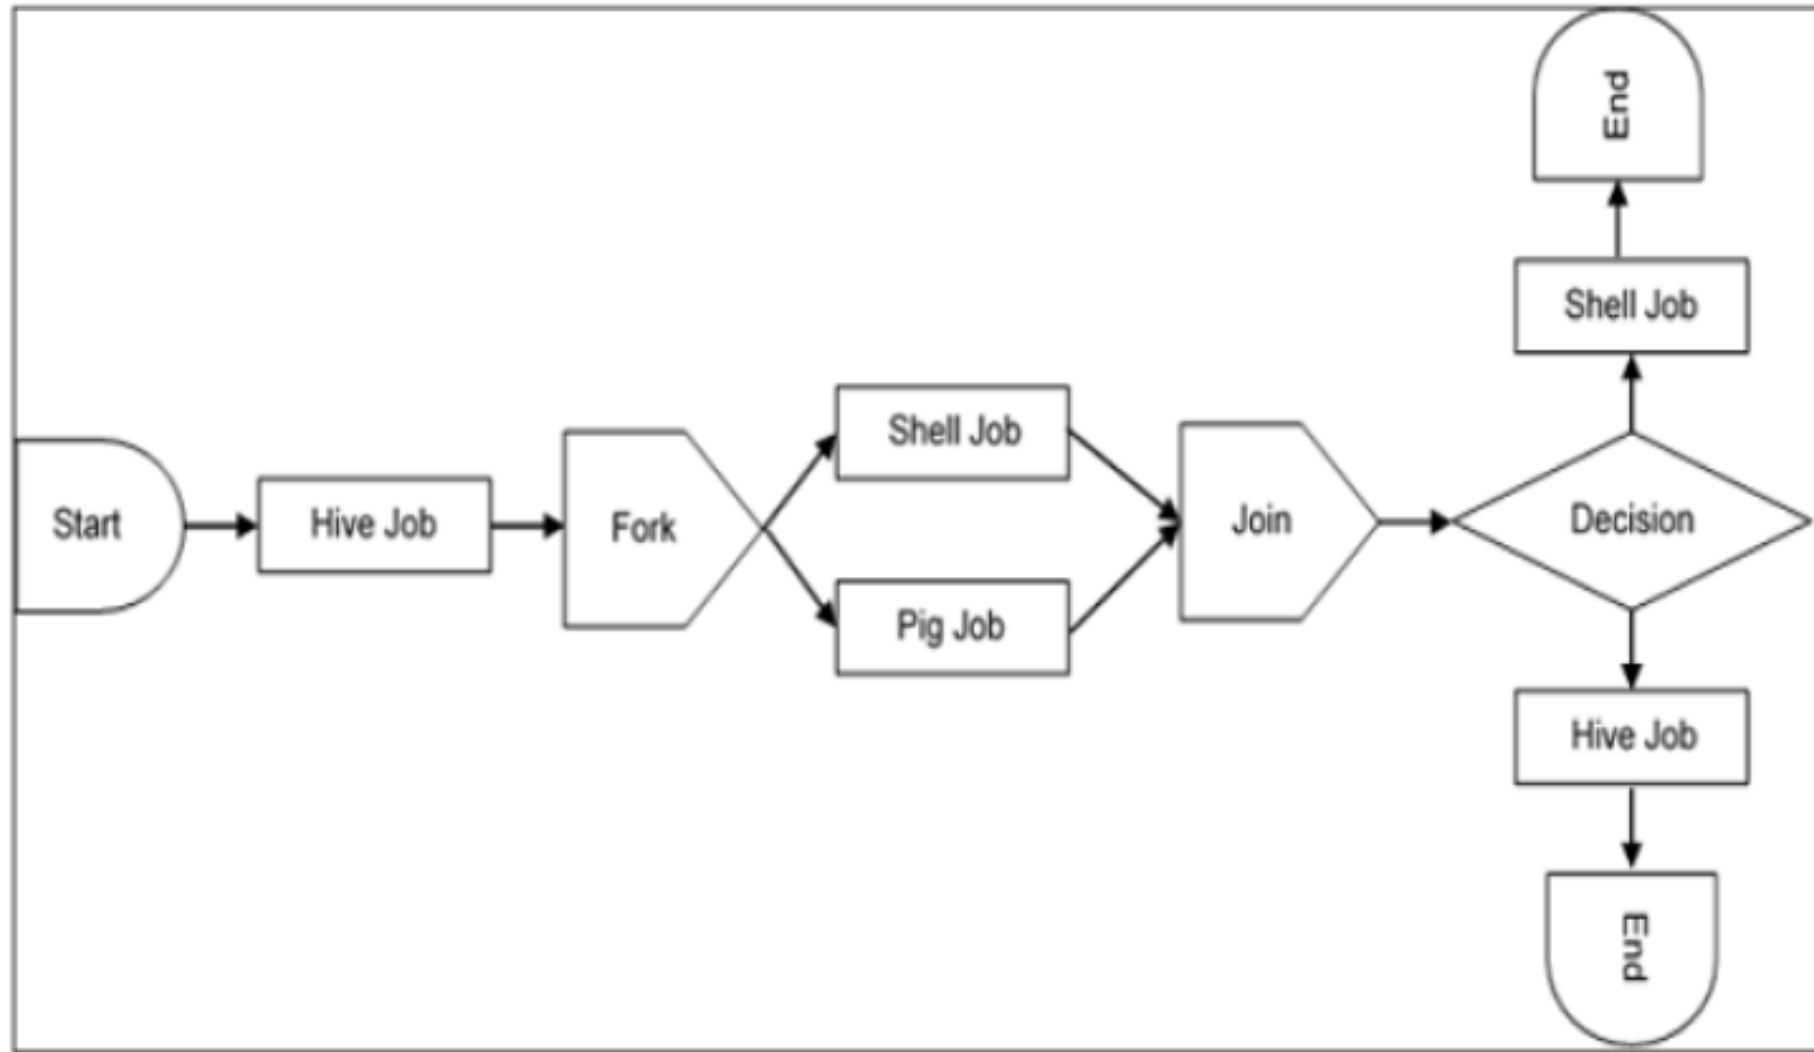
Apache OOZIE: Types of Nodes

- **Start and end nodes:** define the beginning and the end of the workflow. They include optional **fail** nodes.
- **Action Nodes:** are where the actual processing tasks are defined. When an action node finishes, the remote systems notify Oozie and the next node in the workflow is executed. Action nodes can also includes HDFS commands.
- **Fork/Join nodes:** enable parallel execution of tasks in the workflow. The fork node enables two or more tasks to run at the same time. A join node represents a rendezvous point that must wait until all forked tasks complete.
- **Control flow nodes:** enable decisions to be made about the previous task. Control decisions are based on the results of the previous actions. Decision nodes are essentially switch-case statement that use JSP EL (Java Server Pages – Expression Language) that evaluate to either true or false.

Apache Oozie: workflow



Apache OozIE: Complex workflow



THANK
YOU.

* ALL THE BEST *